

Daemons in Ruby

Jose Luis Salas

September 23, 2014

Presentation outline

- 1 Heads up
- 2 Processes
- 3 Daemons
- 4 Pimping our daemon

Let's start

Consider this code

```
loop do
  # Do stuff here
  sleep 1
end
```

Let's start

Consider this code

```
loop do
  # Do stuff here
  sleep 1
end
```

Is this a daemon?

Let's start

Consider this code

```
loop do
  # Do stuff here
  sleep 1
end
```

Is this a daemon? No, it isn't

What's a daemon

Daemons are processes on a system which principal characteristics are:

What's a daemon

Daemons are processes on a system which principal characteristics are:

- Aren't directly controlled by the user

What's a daemon

Daemons are processes on a system which principal characteristics are:

- Aren't directly controlled by the user
- Run in background

What's a daemon

Daemons are processes on a system which principal characteristics are:

- Aren't directly controlled by the user
- Run in background
- Are spawned typically by init manager

What's a daemon

Daemons are processes on a system which principal characteristics are:

- Aren't directly controlled by the user
- Run in background
- Are spawned typically by init manager
- Typically log data into logfiles

What's a daemon

Daemons are processes on a system which principal characteristics are:

- Aren't directly controlled by the user
- Run in background
- Are spawned typically by init manager
- Typically log data into logfiles
- Are detached from terminal

Presentation outline

- 1 Heads up
- 2 Processes
- 3 Daemons
- 4 Pimping our daemon

Process

Process

- What is a process?

Process

- What is a process?
Is an instance of a computer program that is being executed

Process

- What is a process?
Is an instance of a computer program that is being executed
- How processes are created?

Process

- What is a process?
Is an instance of a computer program that is being executed
- How processes are created? fork in UNIX

Process

- What is a process?
Is an instance of a computer program that is being executed
- How processes are created? fork in UNIX, CreateProcess on the rest

Process

- What is a process?
Is an instance of a computer program that is being executed
- How processes are created? fork in UNIX, CreateProcess on the rest
- What is shared between processes?

Process

- What is a process?
Is an instance of a computer program that is being executed
- How processes are created? fork in UNIX, CreateProcess on the rest
- What is shared between processes? Libraries?

Process

- What is a process?
Is an instance of a computer program that is being executed
- How processes are created? fork in UNIX, CreateProcess on the rest
- What is shared between processes? Libraries? Memory?

What is a process consisted of?

What is a process consisted of?

- An image of the executable machine code.

What is a process consisted of?

- An image of the executable machine code.
- Virtual memory which includes the executable code, process-specific data, a call stack, and a heap to hold intermediate computation data generated during run time.

What is a process consisted of?

- An image of the executable machine code.
- Virtual memory which includes the executable code, process-specific data, a call stack, and a heap to hold intermediate computation data generated during run time.
- Operating system descriptors of resources, such as file descriptors.

What is a process consisted of?

- An image of the executable machine code.
- Virtual memory which includes the executable code, process-specific data, a call stack, and a heap to hold intermediate computation data generated during run time.
- Operating system descriptors of resources, such as file descriptors.
- Security attributes, such as the process owner and the process' set of permissions.

What is a process consisted of?

- An image of the executable machine code.
- Virtual memory which includes the executable code, process-specific data, a call stack, and a heap to hold intermediate computation data generated during run time.
- Operating system descriptors of resources, such as file descriptors.
- Security attributes, such as the process owner and the process' set of permissions.
- Processor state, such as the content of registers, physical memory addressing, etc. The state is typically stored in CPU registers when the process is executing, and in memory otherwise.

Executing a program

```
puts("Hello from process #{Process.pid}")  
exec('uname -r')  
puts("Bye from process")
```

Executing a program

```
puts("Hello from process #{Process.pid}")  
exec('uname -r')  
puts("Bye from process")
```

Hello from process 23191
3.14-2-amd64

Executing programs

- Kernel.exec

Executing programs

- `Kernel.exec` : Replaces the current process image
- `Kernel.system`

Executing programs

- `Kernel.exec` : Replaces the current process image
- `Kernel.system` : Executes `cmd` in a subshell
- `Kernel.`` (backticks)

Executing programs

- `Kernel.exec` : Replaces the current process image
- `Kernel.system` : Executes `cmd` in a subshell
- `Kernel.`` (backticks) : Executes `cmd` in a subshell
- `IO.popen`

Executing programs

- `Kernel.exec` : Replaces the current process image
- `Kernel.system` : Executes `cmd` in a subshell
- `Kernel.`` (backticks) : Executes `cmd` in a subshell
- `IO.popen` : Runs command as a subprocess
- `Open3.popen3`
- `Process.spawn`
- `Process.daemon`
- `IO.popen4` (JRuby)

Creating a new process

```
puts("I am process #{Process.pid}")  
fork  
puts("I am process #{Process.pid}")
```

Creating a new process

```
puts("I am process #{Process.pid}")  
fork  
puts("I am process #{Process.pid}")
```

I am process 23823

I am process 23823

I am process 23825

Executing a new program in a new process

```
puts("I am process #{Process.pid}")
pid = fork do
  puts("I am process #{Process.pid}, my
      parent is #{Process.ppid}")
end

puts("I am process #{Process.pid}, I am
    waiting for process #{pid}")
Process.wait(pid)
```

Executing a new program in a new process

```
puts("I am process #{Process.pid}")
pid = fork do
  puts("I am process #{Process.pid}, my
      parent is #{Process.ppid}")
end

puts("I am process #{Process.pid}, I am
    waiting for process #{pid}")
Process.wait(pid)
```

I am process 11135

I am process 11135, I am waiting for process 11137

I am process 11137, my parent is 11135

Trying to become a daemon

```
def print_process_info
  puts("PID: #{Process.pid} PPID:
      #{Process.ppid} SID: #{Process.getsid}
      PGRP: #{Process.getpgrp}")
end

print_process_info
exit if fork
print_process_info
Process.getsid
print_process_info
exit if fork
print_process_info
```

Trying to become a daemon part 2

```
$ ruby process5.rb
```

```
PID: 5764 PPID: 4755 SID: 4755 PGRP: 5764
```

```
PID: 5805 PPID: 5764 SID: 4755 PGRP: 5764
```

```
PID: 5805 PPID: 5764 SID: 4755 PGRP: 5764
```

```
PID: 5808 PPID: 5805 SID: 4755 PGRP: 5764
```

```
$ ps -o ppid , pid , sid , pgid , command
```

```
PPID    PID    SID    PGID  COMMAND
```

```
1      5808   4755   5764  ruby process5.rb
```


Communicating processes with pipes

```
pipe_me_in, pipe_child_out = IO.pipe
pipe_child_in, _ = IO.pipe

fork do
  STDIN.reopen(pipe_child_in)
  STDOUT.reopen(pipe_child_out)
  exec("echo valencia.rb")
end

pipe_child_out.close
puts(pipe_me_in.read)
```

Communicating processes with pipes

```
pipe_me_in, pipe_child_out = IO.pipe
pipe_child_in, _ = IO.pipe

fork do
  STDIN.reopen(pipe_child_in)
  STDOUT.reopen(pipe_child_out)
  exec("echo valencia.rb")
end

pipe_child_out.close
puts(pipe_me_in.read)
```

valencia.rb

Signals

```
trap(:KILL) do
  puts('I am not going to die!')
end
puts('Trying to kill myself')
Process.kill(:KILL, Process.pid)
puts('I should be alive')
```

Signals

```
trap(:KILL) do
  puts('I am not going to die!')
end
puts('Trying to kill myself')
Process.kill(:KILL, Process.pid)
puts('I should be alive')
```

Trying to kill myself
Killed

Presentation outline

- 1 Heads up
- 2 Processes
- 3 Daemons**
- 4 Pimping our daemon

Steps to become a UNIX daemon

Steps to become a UNIX daemon

- Dissociate from the controlling tty

Steps to become a UNIX daemon

- Dissociate from the controlling tty
- Become a session leader

Steps to become a UNIX daemon

- Dissociate from the controlling tty
- Become a session leader
- Become a process group leader

Steps to become a UNIX daemon

- Dissociate from the controlling tty
- Become a session leader
- Become a process group leader
- Execute as a background task by forking and exiting (once or twice), required sometimes to become a session leader.

Steps to become a UNIX daemon

- Dissociate from the controlling tty
- Become a session leader
- Become a process group leader
- Execute as a background task by forking and exiting (once or twice), required sometimes to become a session leader.
- Setting the root directory (/) as the current working directory so that the process does not keep any directory in use.

Steps to become a UNIX daemon

- Dissociate from the controlling tty
- Become a session leader
- Become a process group leader
- Execute as a background task by forking and exiting (once or twice), required sometimes to become a session leader.
- Setting the root directory (/) as the current working directory so that the process does not keep any directory in use.
- Changing the umask to 0 to allow operating system calls to provide their own permission masks and not to depend on the umask of the caller

Steps to become a UNIX daemon

- Dissociate from the controlling tty
- Become a session leader
- Become a process group leader
- Execute as a background task by forking and exiting (once or twice), required sometimes to become a session leader.
- Setting the root directory (/) as the current working directory so that the process does not keep any directory in use.
- Changing the umask to 0 to allow operating system calls to provide their own permission masks and not to depend on the umask of the caller
- Closing all inherited files at the time of execution that are left open by the parent process, including file descriptors 0, 1 and 2.

Steps to become a UNIX daemon

- Dissociate from the controlling tty
- Become a session leader
- Become a process group leader
- Execute as a background task by forking and exiting (once or twice), required sometimes to become a session leader.
- Setting the root directory (/) as the current working directory so that the process does not keep any directory in use.
- Changing the umask to 0 to allow operating system calls to provide their own permission masks and not to depend on the umask of the caller
- Closing all inherited files at the time of execution that are left open by the parent process, including file descriptors 0, 1 and 2.
- Using a logfile, the console, or /dev/null as stdin, stdout, and stderr

Daemonizing in Ruby 1.9

```
Process.daemon
```

Daemonizing in Ruby 1.9

```
Process.daemon
```

- Detach the process from controlling terminal and run in the background as system daemon.
- Unless the argument `nochdir` is true, it changes the current working directory to `/`.
- Unless the argument `noclose` is true, `daemon()` will redirect standard input, standard output and standard error to `/dev/null`.
- Return zero on success, or raise one of `Errno::*`.

proc_daemon

From <https://github.com/ruby/ruby/blob/master/process.c>

```
static VALUE proc_daemon(int argc, VALUE argv) {
    VALUE nochdir, noclose;
    int n;

    rb_secure(2);
    rb_scan_args(argc, argv, "02", &nochdir, &noclose);

    prefork();
    before_fork();
    n = daemon(RTEST(nochdir), RTEST(noclose));
    after_fork();
    if (n < 0) rb_sys_fail("daemon");
    return INT2FIX(n);
}

#define daemon(nochdir, noclose) rb_daemon(nochdir, noclose)
#endif
```

rb_process

From <https://github.com/ruby/ruby/blob/master/process.c>

```
static int rb_daemon(int nochdir, int noclose) {
    int n, err = 0;

    switch (rb_fork(0, 0, 0, Qnil)) {
        case -1:
            rb_sys_fail("daemon");
        case 0:
            break;
        default:
            _exit(EXIT_SUCCESS);
    }

    proc_setsid();

    switch (rb_fork(0, 0, 0, Qnil)) {
        case -1:
            return -1;
        case 0:
            break;
        default:
            _exit(EXIT_SUCCESS);
    }

    if (!nochdir) err = chdir("/");

    if (!noclose && (n = open("/dev/null", O_RDWR, 0)) != -1) {
        (void)dup2(n, 0);
        (void)dup2(n, 1);
        (void)dup2(n, 2);
        if (n > 2) (void)close(n);
    }
    return err;
}
```

rb_fork

From <https://github.com/ruby/ruby/blob/master/process.c>

```
rb_pid_t rb_fork(int *status, int (*chfunc)(void*), void *charg, VALUE fds) {  
    if (chfunc) {  
        struct chfunc_wrapper_t warg;  
        warg.chfunc = chfunc;  
        warg.arg = charg;  
        return rb_fork_err(status, chfunc_wrapper, &warg, fds, NULL, 0);  
    } else {  
        return rb_fork_err(status, NULL, NULL, fds, NULL, 0);  
    }  
}
```

rb_fork_err

From <https://github.com/ruby/ruby/blob/master/process.c>

```
rb_pid_t rb_fork_err(int *status, int (*chfunc)(void*, char *, size_t),
    void *charg, VALUE fds, char *errmsg, size_t errmsg_buflen) {
    rb_pid_t pid;
    int err, state = 0;

#define prefork() (define\
    rb_io_flush(rb_stdout), \
    rb_io_flush(rb_stderr)rb_stderr\
    )
    prefork();
    // STUFF
    for (; before_fork(), (pid = fork()) < 0; prefork()) {
        after_fork();
        // STUFF
    }
    if (!pid) {
        forked_child = 1;
        if (chfunc) {
            if (!(*chfunc)(charg, errmsg, errmsg_buflen)) _exit(EXIT_SUCCESS);
        }
    }
    if (EXIT_SUCCESS == 127
        _exit(EXIT_FAILURE);
    }
    else
        _exit(127);
    }
    after_fork();
    return pid;
}
```

proc_setsid

From <https://github.com/ruby/ruby/blob/master/process.c>

```
static VALUE proc_setsid(void) {  
    rb_pid_t pid;  
  
    rb_secure(2);  
    pid = setsid();  
    if (pid < 0) rb_sys_fail(0);  
    return PIDT2NUM(pid);  
}  
  
#define setsid() ruby_setuid()
```

ruby_setsid.c

From <https://github.com/ruby/ruby/blob/master/process.c>

```
static rb_pid_t ruby_setsid(void) {
    rb_pid_t pid;
    int ret;

    pid = getpid();
    #if defined(SETPGRP_VOID)
        ret = setpgrp();
        /* If 'pid_t setpgrp(void)' is equivalent to setsid(),
         * 'ret' will be the same value as 'pid', and following open() will fail.
         * In Linux, 'int setpgrp(void)' is equivalent to setpgid(0, 0). */
    #else
        ret = setpgrp(0, pid);
    #endif
    if (ret == -1) return -1;

    if ((fd = open("/dev/tty", O_RDWR)) >= 0) {
        ioctl(fd, TIOCNOTTY, NULL);
        close(fd);
    }
    return pid;
}
#endif
```

proc_setpgrp.c

From <https://github.com/ruby/ruby/blob/master/process.c>

```
static VALUE proc_setpgrp(void) {  
  rb_secure(2);  
  /* check for posix setpgid() first; this matches the posix */  
  /* getpgrp() above. It appears that configure will set SETPGRP_VOID */  
  /* even though setpgrp(0,0) would be preferred. The posix call avoids */  
  /* this confusion. */  
#ifndef HAVE_SETPGID  
  if (setpgid(0,0) < 0) rb_sys_fail(0);  
#elif defined(HAVE_SETPGRP) && defined(SETPGRP_VOID)  
  if (setpgrp() < 0) rb_sys_fail(0);  
#endif  
  return INT2FIX(0);  
}  
#endif
```

How to daemonize a program in Ruby

```
def daemonize_app
  exit if fork
  Process.setsid
  exit if fork
  Dir.chdir("/")
  STDIN.reopen("/dev/null")
  STDOUT.reopen("/dev/null", "a")
  STDERR.reopen("/dev/null", "a")
end
```


Checking it live

```
$ ls -l /proc/2206/fd
total 0
```

```
lrwx----- 1 selu selu 64 Sep 21 16:52 0 -> /dev/null
lrwx----- 1 selu selu 64 Sep 21 16:52 1 -> /dev/null
lrwx----- 1 selu selu 64 Sep 21 16:52 2 -> /dev/null
lr-x----- 1 selu selu 64 Sep 21 16:52 3 -> pipe:[595663]
l-wx----- 1 selu selu 64 Sep 21 16:52 4 -> pipe:[595663]
lr-x----- 1 selu selu 64 Sep 21 16:52 5 -> pipe:[595664]
l-wx----- 1 selu selu 64 Sep 21 16:52 6 -> pipe:[595664]
```

```
$ ps xf -o pid,ppid,pgid,sid,command | grep -e [r]uby -e [P]ID
PID   PPID   PGID   SID COMMAND
2206   1      2203   2203 ruby -e Process.daemon; sleep 3600
```

Daemonize in daemons gem

From <https://github.com/ghazel/daemons/blob/master/lib/daemons/daemonize.rb>

```
def daemonize(logfile_name = nil, app_name = nil)
  srand
  safe_fork and exit

  unless sess_id = Process.setsid
    raise Daemons::RuntimeError.new('cannot detach from controlling
                                     terminal')
  end

  # Prevent the possibility of acquiring a controlling terminal
  trap 'SIGHUP', 'IGNORE'
  exit if pid = safe_fork

  $0 = app_name if app_name

  Dir.chdir "/"
  File.umask 0000

  ObjectSpace.each_object(IO) do |io|
    unless [STDIN, STDOUT, STDERR].include?(io)
      begin
        io.close unless io.closed?
      rescue ::Exception
      end
    end
  end

  redirect_io(logfile_name)
  return sess_id
end
```

Presentation outline

- 1 Heads up
- 2 Processes
- 3 Daemons
- 4 Pimping our daemon

TCP preforking daemon

```
require 'socket'
Process.daemon
socket = TCPServer.open('0.0.0.0', 8080)
wpids = []

5.times do
  wpids << fork do
    loop do
      connection = socket.accept
      connection.puts "Hello from #{Process.pid}"
      connection.close
    end
  end
end

[:INT, :QUIT].each do |signal|
  Signal.trap(signal) do
    wpids.each { |wpid| Process.kill(signal, wpid) }
  end
end
Process.waitall
```

```
$ nc localhost 8080;nc localhost 8080
Hello from 8349
Hello from 8361
```

Improving signals

```
def handle_signal(signal)
  puts("Received #{signal}")
end

self_read, self_write = IO.pipe

%w(INT TERM).each do |sig|
  trap(sig) { self_write.puts(sig) }
end

while readable_io = IO.select([self_read])
  signal = readable_io.first[0].gets.strip
  handle_signal(signal)
end
```

Our robust daemon

```
Process.daemon
socket = TCPServer.open('0.0.0.0', 8080)
wpids = []

5.times do
  wpids << fork do
    run = true
    [:INT, :QUIT].each do |signal|
      Signal.trap(signal, lambda { run = false })
    end
    while run
      connection = socket.accept
      connection.puts("Hello from #{Process.pid}")
      sleep 10
      connection.puts("Goodbye from #{Process.pid}")
      connection.close
    end
  end
end

[:INT, :QUIT].each do |signal|
  Signal.trap(signal) do
    wpids.each { |wpid| Process.kill(signal, wpid) }
  end
end
Process.waitall
```

```
$ nc localhost 8080; killall -INT ruby
Hello from 8328
Goodbye from 8328
```

Our robust daemon with logging

```
Process.daemon
socket = TCPServer.open('0.0.0.0', 8080)
wpids = []

logger = Logger.new('/tmp/logger.log')
logger.level = Logger::INFO

5.times do
  wpids << fork do
    run = true
    [:INT, :QUIT].each do |signal|
      Signal.trap(signal, lambda { run = false })
    end
    while run
      connection = socket.accept
      connection.puts("Hello from #{Process.pid}")
      logger.info("Connection in #{Process.pid}")
      sleep 10
      connection.puts("Goodbye from #{Process.pid}")
      connection.close
    end
  end
end

[:INT, :QUIT].each do |signal|
  Signal.trap(signal) do
    wpids.each { |wpid| Process.kill(signal, wpid) }
  end
end
Process.waitall
logger.close
```

```
$ tail -f /tmp/logger.log
# Logfile created on 2014-09-21 19:46:25 +0200 by logger.rb/v1.2.7
I, [2014-09-21T19:47:12.696949 #9452] INFO — : Connection in 9452
I, [2014-09-21T19:47:22.702008 #9449] INFO — : Connection in 9449
```

Logs rotation with a signal

From https://github.com/kennethkalmer/daemon-kit/blob/master/lib/daemon_kit/application.rb

```
configuration.trap("HUP") {
  DaemonKit::Application.reopen_logs
}

module DaemonKit
  class Application
    def reopen_logs
      nr = 0
      append_flags = File::WRONLY | File::APPEND
      DaemonKit.logger.info "Rotating logs" if DaemonKit.logger

      ObjectSpace.each_object(File) do |fp|
        next if fp.closed?
        next unless (fp.sync && fp.path[0..0] == "/")
        next unless (fp.fcntl(Fcntl::F_GETFL) & append_flags) == append_flags

        begin
          a, b = fp.stat, File.stat(fp.path)
          next if a.ino == b.ino && a.dev == b.dev
        rescue Errno::ENOENT
        end

        open_arg = 'a'
        if fp.respond_to?(:external_encoding) && enc = fp.external_encoding
          open_arg << ":{enc.to_s}"
          enc = fp.internal_encoding and open_arg << ":{enc.to_s}"
        end
        DaemonKit.logger.info "Rotating path: #{fp.path}" if DaemonKit.logger
        fp.reopen(fp.path, open_arg)
        fp.sync = true
        nr += 1
      end # each_object
      nr
    end
  end
end
```


The end

The end

Further info:

- http://codeincomplete.com/posts/2014/9/15/ruby_daemons/
- Working with unix processes by Jesse Storimer
- lib/unicorn/launcher.rb in Unicorn gem
- daemons, dante and daemon-kit gems
- Linux System Programming: Talking Directly to the Kernel and C Library
- Understanding the Linux Kernel

The end

Further info:

- http://codeincomplete.com/posts/2014/9/15/ruby_daemons/
- Working with unix processes by Jesse Storimer
- lib/unicorn/launcher.rb in Unicorn gem
- daemons, dante and daemon-kit gems
- Linux System Programming: Talking Directly to the Kernel and C Library
- Understanding the Linux Kernel

Questions?

The end

Further info:

- http://codeincomplete.com/posts/2014/9/15/ruby_daemons/
- Working with unix processes by Jesse Storimer
- lib/unicorn/launcher.rb in Unicorn gem
- daemons, dante and daemon-kit gems
- Linux System Programming: Talking Directly to the Kernel and C Library
- Understanding the Linux Kernel

Questions? Thanks!